Action	Code	Explantion
Mount Drive	`#mounting drive` `from google.colab import drive` `drive.mount('/content/drive')`	
Import Panda	`import pandas as pd`	
Upload CSV file	`df1=pd.read_csv("/content/drive/MyDrive/wk10_pandas_learner_vs1902/Resources/Flight Delays/Flights_Delay_Cause_2019-2020.csv")` `df2=pd.read_csv("/content/drive/MyDrive/wk10_pandas_learner_vs1902/Resources/Flight Delays/Flights_Delay_Cause_2020-2021.csv")` `df3=pd.read_csv("/content/drive/MyDrive/wk10_pandas_learner_vs1902/Resources/Flight Delays/Flights_Delay_Cause_2021-2022.csv")`	
To see 1st 5 & last 5 lines	`df1.head()` `df1tail()`	
To combine multiple CSV file in 1 dataset	`#use function .concat()` `df= pd.concat([df1, df2, df3], ignore_index=True)` `#setting the parameterignore_index=True will reasign new indexes to the whole dataset, the intial indexes of each document will be ignored`	
Explore wholedataset	`df.head(10)`	
Explore total rows &Columns	`df.info()` `<class 'pandas.core.frame.DataFrame'>` `RangeIndex: 42461 entries, 0 to 42460` `Data columns (total 16 columns):`	
Attribution step	`df.shape` (42461, 16)	

Datatype	`df.dtypes` year int64 month int64 carrier object carrier_name object airport object airport_name object arr_flights float64 arr_del15 float64 arr_cancelled float64 arr_diverted float64 arr_delay float64 carrier_delay float64 weather_delay float64 nas_delay float64 security_delay float64 delays float64 dtype: object	
Describe	`df.decribe()`	
Convert data into a date datatype	`.todatetime()`	
Copy original dataset to avoid errors	`df_copy= df.copy(deep=True)`	This is our new copied database name from now on.
Add “Date” column	the "month" and "year" from our dataset. We can concatenate our 2 columns month and year using the "+" sign: `df_copy['date']= df_copy['month']+df_copy['year']` `df_copy.head()`	Error spotted in date column

	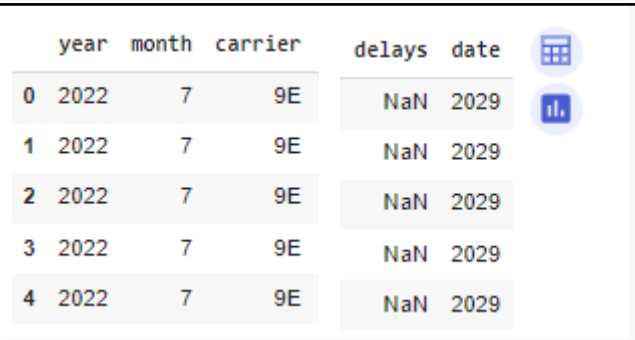 year month carrier delays date 0 2022 7 9E NaN 2029 1 2022 7 9E NaN 2029 2 2022 7 9E NaN 2029 3 2022 7 9E NaN 2029 4 2022 7 9E NaN 2029	
Change the month & year datatype to string	.astype() `#First, let's convert the month and year into a string datatype so we can concatenate them afterall:` `#We need to use the .astype() methodf to convert data in Pandas` `df_copy=df_copy.astype({"month": str, "year":str, "date":str})` `df_copy.info()` `<class 'pandas.core.frame.DataFrame'>` `RangeIndex: 42461 entries, 0 to 42460` `Data columns (total 17 columns):` ` # Column Non-Null Count Dtype` `--- ------ -------------- -----` ` 0 year 42461 non-null object` ` 1 month 42461 non-null object` **OR** `#alternative method to convert datatypes in Pandas:` `df_copy['month'].astype(str)` `df_copy['year'].astype(str)` `df_copy['date'].astype(str)`	Datatype changed to object in pandas

	df_copy['date']= df_copy['year']+"-"+df_copy['month'] df_copy.head() 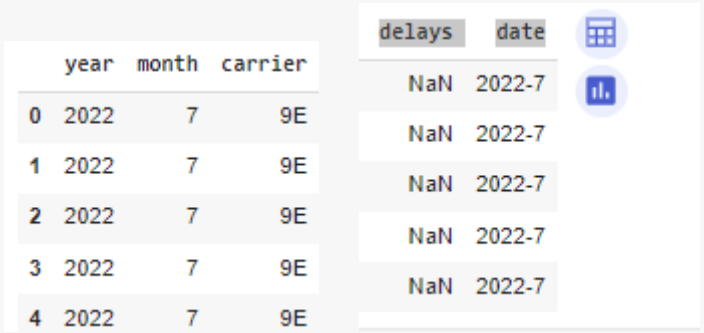	
Transform the date into datetime datatype:	to_datetime() df_copy['date']= pd.to_datetime(df_copy['date']) df_copy.info() 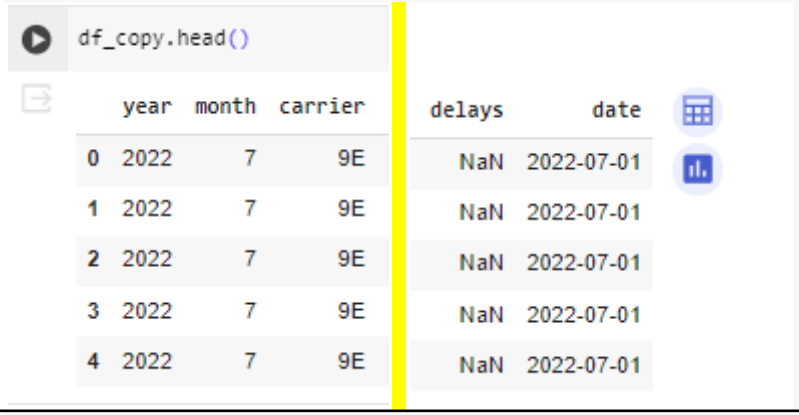	
Correct date format	df_copy['date']= pd.to_datetime(df_copy['date']).dt.to_period('M')	

	year	month	carrier
0	2022	7	9E
1	2022	7	9E
2	2022	7	9E
3	2022	7	9E
4	2022	7	9E

delays	date
NaN	2022-7
NaN	2022-7
NaN	2022-7
NaN	2022-7
NaN	2022-7

	year	month	carrier
0	2022	7	9E
1	2022	7	9E
2	2022	7	9E
3	2022	7	9E
4	2022	7	9E

delays	date
NaN	2022-07-01
NaN	2022-07-01
NaN	2022-07-01
NaN	2022-07-01
NaN	2022-07-01

	16 date 42461 non-null period[M] dtypes: float64(10), object(6), period[M](1) memory usage: 5.5+ MB The 'date' column was assigned the 'period' datatype which is another date fromat	
Save added “Date” column in our database	df_copy.to_csv('flights_dataset.csv') and mount the drive from google.colab import drive drive.mount('/content/drive')	To save in different location from google.colab import drive; drive.mount('/content/drive'). Then, use a path within your Drive, like /content/drive/MyDrive/data/flights_dataset.csv.
Clear unwanted columns	df_copy=df_copy.drop(['month', 'year'], axis=1) df_copy.head()	
3. DATA CLEANING		
Find missing values	isnull() df_copy.isnull().sum()	

	Let's check how many missing values our dataset has using the isnull() functio [31] `df_copy.isnull().sum()` carrier 0 carrier_name 16 airport 0 airport_name 0 arr_flights 107 arr_del15 278 arr_cancelled 107 arr_diverted 107 arr_delay 107 carrier_delay 107 weather_delay 107 nas_delay 107 security_delay 107 delays 42461 date 0 dtype: int64	
Drop column	`df_copy= df_copy.dropna(subset=['carrier_name'])` `df_copy.isnull().sum()`	
Get column names	`Df_copy.columns` Index(['carrier', 'carrier_name', 'airport', 'airport_name', 'arr_flights', 'arr_del15', 'arr_cancelled', 'arr_diverted', 'arr_delay', 'carrier_delay', 'weather_delay', 'nas_delay', 'security_delay', 'delays', 'date'], dtype='object')	
Built in box option	`.box_plot` `df_copy.boxplot('weather_delay')` `#We create a function that will return a box plot:`	`# Identify and remove the row with the outlier`

Name the chart X Axis	def box_plot(flights): return df_copy.boxplot(flights) #We call our function passing a data variable as argument: box_plot('nas_delay')	outlier_index= df_copy[df_copy['nas_delay']>175000].index df_cleaned= df_copy.drop(outlier_index) Int64Index([363], dtype='int64') #row number
	df_cleaned.boxplot('nas_delay')	
Python marplotlib visualisation	Subplots() import matplotlib.pyplot as plt figure, ax = plt.subplots(nrows=2, ncols=1) # subplot layer has 2 rows an 1 column which means 2 figures df_cleaned.hist('weather_delay',ax=ax[0], bins=25, color="purple") df_cleaned.boxplot("weather_delay",ax=ax[1], color="blue")	<Axes: > 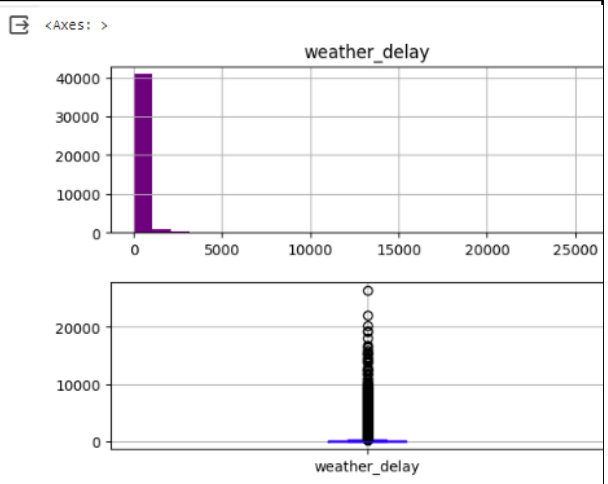
	def subplot_function(data1, data2): # I added one parameter per plot figure, ax = plt.subplots(nrows=2, ncols=1) # Get statistics min_val= df_cleaned[data1].min() max_val= df_cleaned[data1].max() mean_val= df_cleaned[data1].mean() med_val= df_cleaned[data1].median()	def subplot_function(data1, data2): # Define a function to create a multi-plot figure """ Creates a figure with two subplots: a histogram for data1 and a boxplot for data2. Adds reference lines for min, mean

	`` # Add lines for the min, mean and median, and max ax[0].axvline(x=min_val, color = 'gray', linestyle='dashed', linewidth = 2) ax[0].axvline(x=mean_val, color = 'cyan', linestyle='dashed', linewidth = 2) ax[0].axvline(x=med_val, color = 'red', linestyle='dashed', linewidth = 2) ax[0].axvline(x=max_val, color = 'blue', linestyle='dashed', linewidth = 2) #Change labels- 1st plot ax[0].set_ylabel('Frequency') df_cleaned.hist(data1,ax=ax[0], bins=10, color="purple") #Change labels- 2nd plot ax[1].set_xlabel('') ax[1].set_ylabel('Frequency') df_cleaned.boxplot(data2,ax=ax[1], color="blue") ```	``` , median, and max values in the histog ram. Args: data1 (str): Name of the colum n in df_cleaned to be used for the his togram. data2 (str): Name of the colum n in df_cleaned to be used for the box plot. """ figure, ax = plt.subplots(nrows=2, ncols=1) # Create a figure with 2 ro ws and 1 column of subplots # Get statistics for data1 to add reference lines min_val = df_cleaned[data1].min() max_val = df_cleaned[data1].max() mean_val = df_cleaned[data1].mean( ) med_val = df_cleaned[data1].median () # Add vertical reference lines for min, mean, median, and max in the fir st subplot ax[0].axvline(x=min_val, color='gr ay', linestyle='dashed', linewidth=2, label='Min') ax[0].axvline(x=mean_val, color='c yan', linestyle='dashed', linewidth=2, label='Mean') ax[0].axvline(x=med_val, color='re d', linestyle='dashed', linewidth=2, l abel='Median') ax[0].axvline(x=max_val, color='bl ue', linestyle='dashed', linewidth=2, label='Max') ax[0].legend() # Add a legend to ```

		explain the reference lines # Create histogram for data1 in th e first subplot ax[0].set_ylabel('Frequency') # S et y-axis label df_cleaned.hist(data1, ax=ax[0], b ins=10, color="purple") # Create boxplot for data2 in the second subplot ax[1].set_xlabel('') # Remove x-a xis label for cleaner presentation ax[1].set_ylabel('Frequency') # S et y-axis label df_cleaned.boxplot(data2, ax=ax[1] , color="blue")
	subplot_function('nas_delay',"nas_delay")	
Add legend in chart	# Same plot with Labels def subplot_function(data1, data2): # Define a function to create a multi-plot figure """ Creates a figure with two subplots: a histogram for data1 and a boxplot for data2. Adds reference lines for min, mean, median, and max values in the histogram. Args: data1 (str): Name of the column in df_cleaned to be used for the histogram. data2 (str): Name of the column in df_cleaned to be used for the boxplot. """	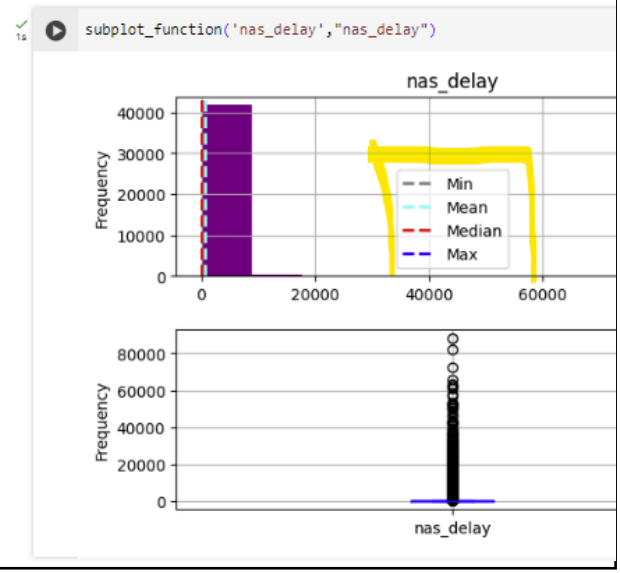

```
    figure, ax = plt.subplots(nrows=2, ncols=1)  # Create a figure with 2 rows and 1 column of subplots

    # Get statistics for data1 to add reference lines
    min_val = df_cleaned[data1].min()
    max_val = df_cleaned[data1].max()
    mean_val = df_cleaned[data1].mean()
    med_val = df_cleaned[data1].median()

    # Add vertical reference lines for min, mean, median, and max in the first subplot
    ax[0].axvline(x=min_val, color='gray', linestyle='dashed', linewidth=2, label='Min')
    ax[0].axvline(x=mean_val, color='cyan', linestyle='dashed', linewidth=2, label='Mean')
    ax[0].axvline(x=med_val, color='red', linestyle='dashed', linewidth=2, label='Median')
    ax[0].axvline(x=max_val, color='blue', linestyle='dashed', linewidth=2, label='Max')
    ax[0].legend()  # Add a legend to explain the reference lines

    # Create histogram for data1 in the first subplot
    ax[0].set_ylabel('Frequency')  # Set y-axis label
    df_cleaned.hist(data1, ax=ax[0], bins=10, color="purple")

    # Create boxplot for data2 in the second subplot
    ax[1].set_xlabel('')  # Remove x-axis label for cleaner presentation
    ax[1].set_ylabel('Frequency')  # Set y-axis label
    df_cleaned.boxplot(data2, ax=ax[1], color="blue")
```

4 Data Analysis

Q1-Which Airlines have the most delayed flights?

We can group the airline data and calculate the average of delayed flights per airline

```
df_delays=df_cleaned.groupby(by="carrier_name", as_index=False)
#using the argument as_index=False in the groupby function will allow us to return a pandas dataframe,
otherwise a pandas series will be returned.
```

	.size() df_delays.size() 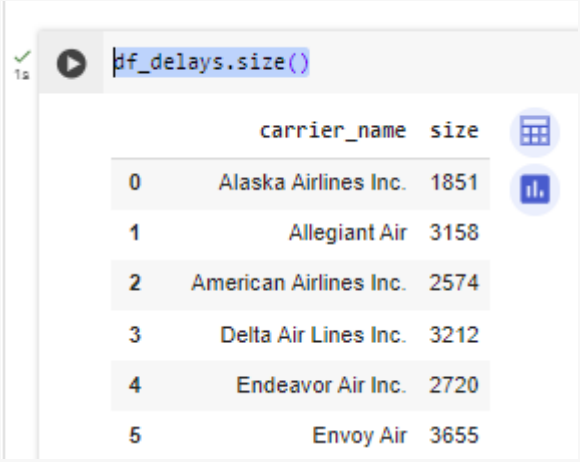 	we can see how many airlines has been grouped with the function .size():
Mean	df_delays=df_delays["arr_del15"].mean() 40.9383…	df_delays # here we get the top 5 of airlines with the most delays
	plot the data using the Matplotlib and Seaborn visualisation libraries import matplotlib.pyplot as plt import seaborn as sns	

	carrier_name	size
0	Alaska Airlines Inc.	1851
1	Allegiant Air	3158
2	American Airlines Inc.	2574
3	Delta Air Lines Inc.	3212
4	Endeavor Air Inc.	2720
5	Envoy Air	3655